

# Technical Report - Bytecode VM

## Architecture of the VM

The virtual machine follows a stack-based execution model. All computations are performed using an operand stack of 32-bit signed integers. A separate return stack is used to manage function calls and returns. The program counter (PC) tracks the current bytecode instruction. A global memory array provides persistent storage accessible via LOAD and STORE instructions.

## Instruction Dispatch Strategy

The VM uses a fetch-decode-execute loop. At each step, one opcode byte is fetched and dispatched using a switch-case construct.

Instructions may be fixed-length or variable-length depending on whether they include operands. The PC is updated accordingly to ensure correct sequential execution.

## Call Frames and Return Mechanism

Function calls are implemented using CALL and RET instructions. CALL saves the return address on a dedicated return stack before jumping to the function entry point. RET restores the saved address.

This design supports nested calls but does not implement local variables or activation records.

## Limitations and Possible Enhancements

Possible future enhancements include support for local stack frames, function arguments, debugging and tracing facilities, and performance optimizations such

- as instruction caching or just-in-time compilation.